

# COMS E6998: Microservices and Cloud Applications

*Lecture 9: CAP Theorem, Database Models, XYZ-Scaling, 12 Factor Apps, Graph DB, Redis*

Dr. Donald F. Ferguson  
dff9@columbia.edu

© Donald F. Ferguson, 2017. All rights reserved.

# Comments Questions

Databases

Consistency

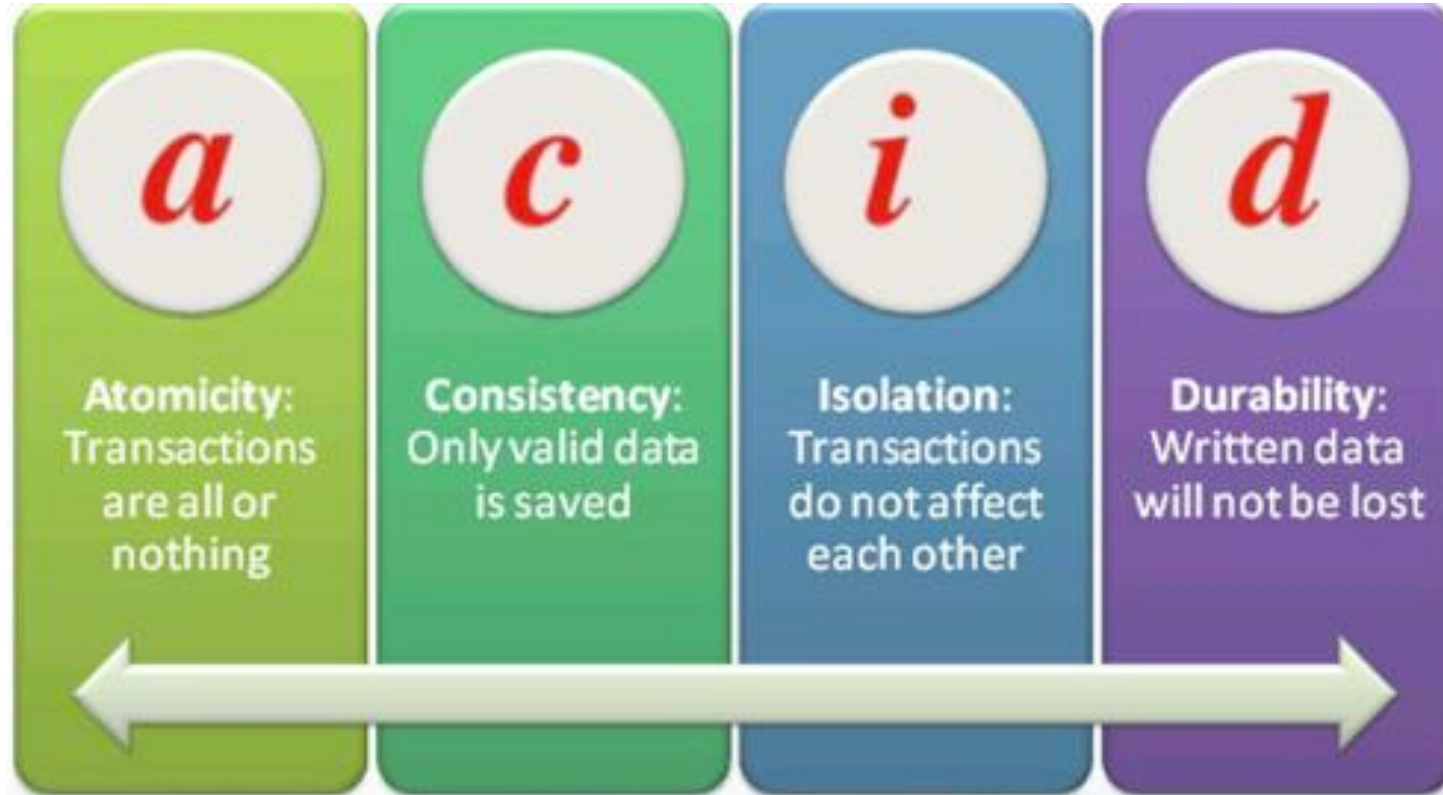
CAP

Theorem

Database Models

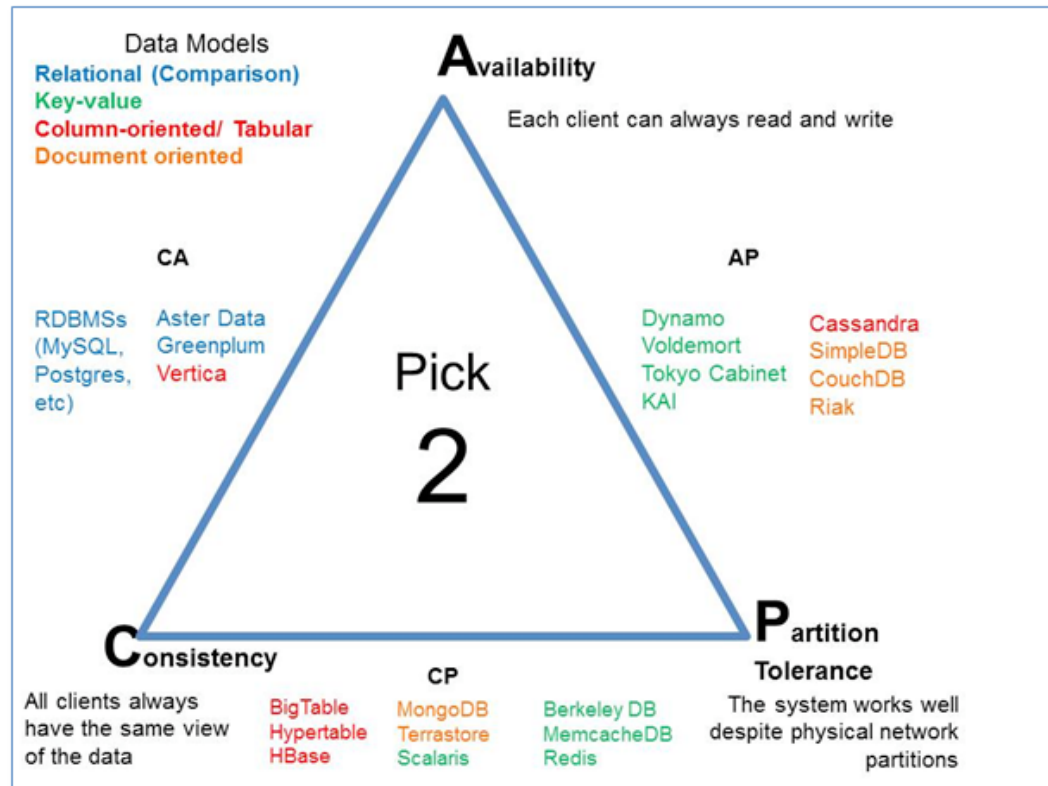
# ACID Databases

(<https://www.linkedin.com/pulse/what-acid-properties-database-aseem-jain/>)



# CAP Theorem

- **Consistency**  
Every read receives the most recent write or an error.
- **Availability**  
Every request receives a (non-error) response – without guarantee that it contains the most recent write.
- **Partition Tolerance**  
The system continues to operate despite an arbitrary number of messages being dropped (or delayed) by the network between nodes



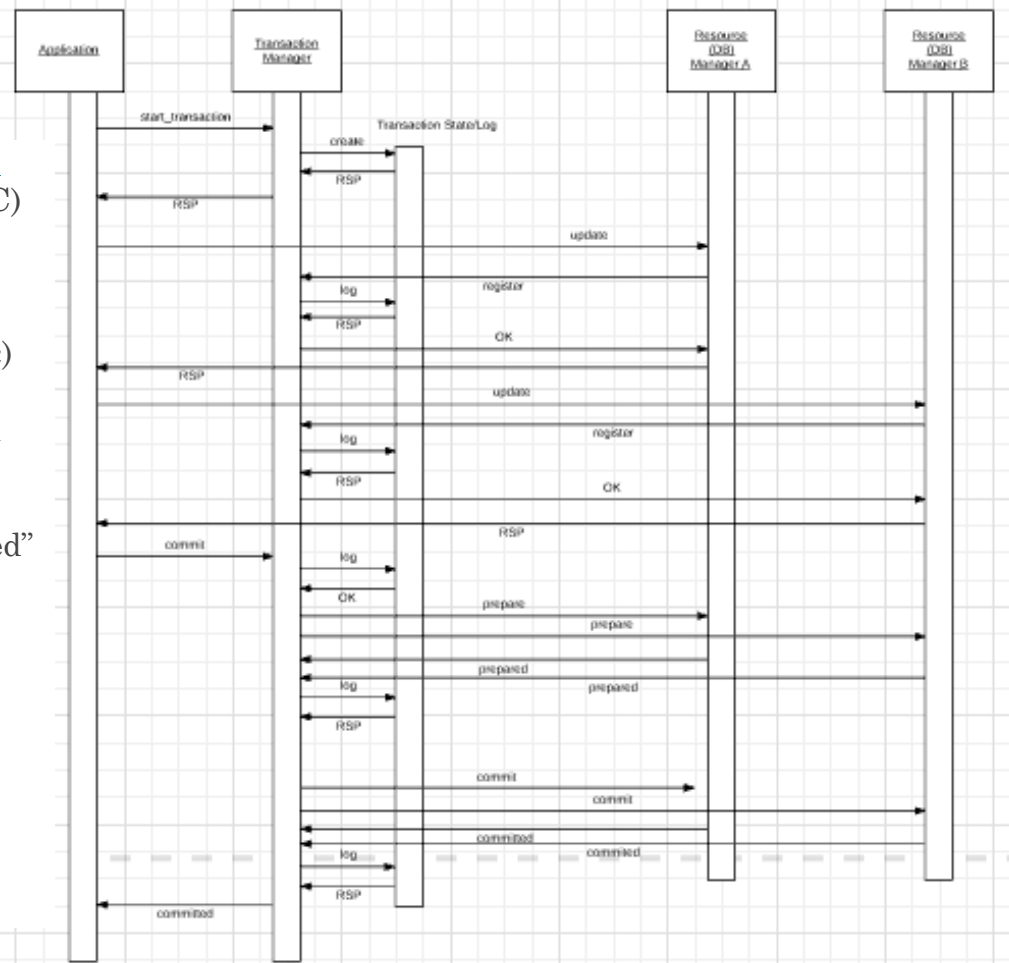
# Two Phase Commit

“In [transaction processing](#), [databases](#), and [computer networking](#), the **two-phase commit protocol (2PC)** is a type of [atomic commitment protocol](#) (ACP). It is a [distributed algorithm](#) that coordinates all the processes that participate in a [distributed atomic transaction](#) on whether to [commit](#) or *abort (roll back)* the transaction (it is a specialized type of [consensus](#) protocol). The protocol achieves its goal even in many cases of temporary system failure (involving either process, network node, communication, etc. failures), and is thus widely used” ([www.wikipedia.org](http://www.wikipedia.org))

Note: This sequence diagram omits some steps.

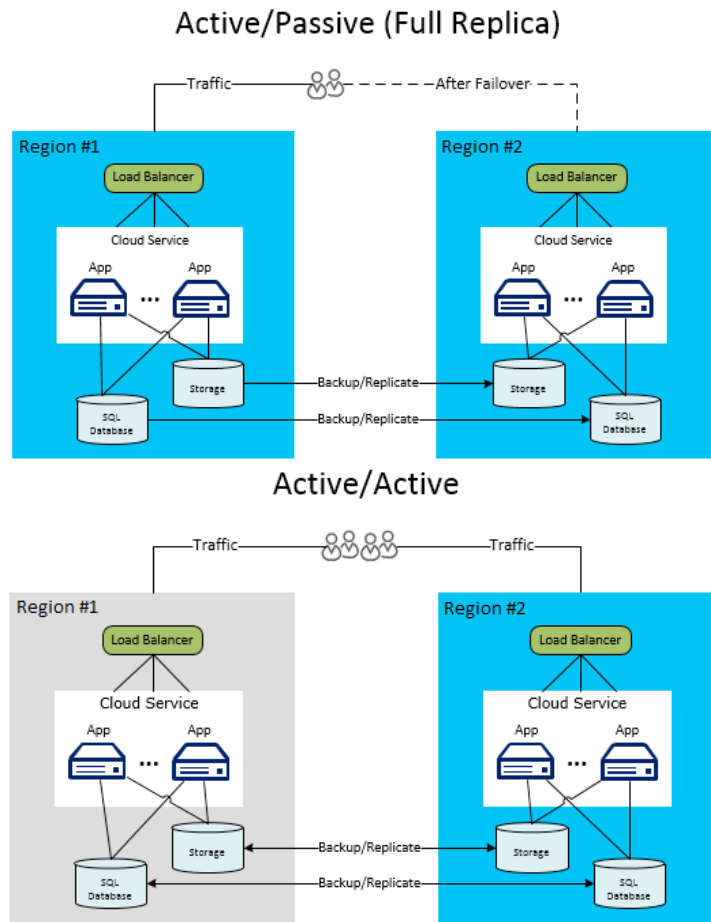
Uses:

- ACID, especially consistency, for data in different databases/services.
- Modified approach is required to strictly achieve consistency and availability for replicated data,



# Availability and Replication

- There are two basic patterns
  - Active/Passive
    - All requests go to *master* during normal processing.
    - Updates are transactionally queued for processing at passive backup.
    - Failure of *master*
      - Routes subsequent requests to *backup*.
      - Backup must process and commit updates before accepting requests.
  - Active/Active
    - Both environments process requests.
    - Some form of distributed transaction commit required to synchronize updates on both copies.
- Multi-system communication to guarantee consistency is the foundation for tradeoffs in CAP.
  - The system can be CAP if and only iff
  - There are never any partitions or system failures
  - Which is unrealistic in cloud/Internet systems.



# Core Data Models

(<http://blog.nahurst.com/visual-guide-to-nosql-systems>)

In addition to CAP configurations, another significant way data management systems vary is by the data model they use: relational, key-value, column-oriented, or document-oriented (there are [others](#), but these are the main ones).

- **Relational** systems are the databases we've been using for a while now. RDBMSs and systems that support ACIDity and joins are considered relational.
- **Key-value** systems basically support get, put, and delete operations based on a primary key.
- **Column-oriented** systems still use tables but have no joins (joins must be handled within your application). Obviously, they store data by column as opposed to traditional row-oriented databases. This makes aggregations much easier.
- **Document-oriented** systems store structured "documents" such as JSON or XML but have no joins (joins must be handled within your application). It's very easy to map data from object-oriented software to these systems.



# Core Data Models

(<http://blog.nahurst.com/visual-guide-to-nosql-systems>)

In addition to CAP configurations, another significant way data management systems vary is in their orientation.

There are also graph databases, which we will cover.

- **Document-oriented** systems store structured "documents" such as JSON or XML but have no joins (joins must be handled within your application). It's very easy to map data from object-oriented software to these systems.

# Database Models

**Available, Partition-Tolerant (AP) Systems** achieve "eventual consistency" through replication and verification. Examples of AP systems include:

- **DynamoDB (key-value)**
- Voldemort (key-value)
- Tokyo Cabinet (key-value)
- KAI (key-value)
- Cassandra (column-oriented/tabular)
- CouchDB (document-oriented)
- SimpleDB (document-oriented)
- Riak (document-oriented)

**Consistent, Available (CA) Systems** have trouble with partitions and typically deal with it with replication.

Examples of CA systems include:

- Traditional RDBMSs like Postgres, MySQL, (relational)
- Vertica (column-oriented)
- Aster Data (relational)
- Greenplum (relational)

**Consistent, Partition-Tolerant (CP) Systems** have trouble with availability while keeping data consistent across partitioned nodes. Examples of CP systems include:

- BigTable (column-oriented/tabular)
- Hypertable (column-oriented/tabular)
- HBase (column-oriented/tabular)
- MongoDB (document-oriented)
- Terrastore (document-oriented)
- **Redis (key-value)**
- Scalaris (key-value)
- MemcacheDB (key-value)
- Berkeley DB (key-value)

<http://blog.nahurst.com/visual-guide-to-nosql-systems>

# Data

Available

"eventual"

verification

- Dyr

- Vol

- Tok

- KA

- Cas

- Cou

- Sim

- Rial

- Key-Value versus Document can be a little confusing.
- Many, if not most, Key-Value stores hold documents.
- The difference is
  - Not in what the data store holds
  - But how you can query and find it.
- Key-Value store finds data by key values explicitly associated with a document.
- Document stores allow queries on the fields/values in the document, similar to a relational WHERE on columns.

<http://blog.nahurst.com/visual-guide-to-nosql-systems>

# Some Queries

- MongoDB Query
  - `db.inventory.find( { status: "A", $or: [ { qty: { $lt: 30 } }, { item: /^p/ } ] } )`
  - Finds all documents in the collection where the status equals "A" **and** *either* qty is less than (\$lt) 30 *or* item starts with the character "p"
- Redis
  - `SADD product:category:1 5 10` // Put product IDs 5 and 10 in category 1.
  - `SMEMBERS product:category:2` // Get the IDs of all products in category 2.
- DynamoDB

Dynamo DB

  - `var params =`  
    `{ TableName : "Movies",`  
    `KeyConditionExpression: "#yr = :yyyy",`  
    `ExpressionAttributeNames: { "#yr": "year" },`  
    `ExpressionAttributeValues: { ":yyyy": 1985 }`  
    `};`

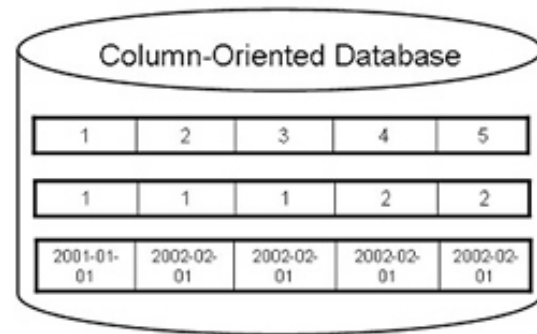
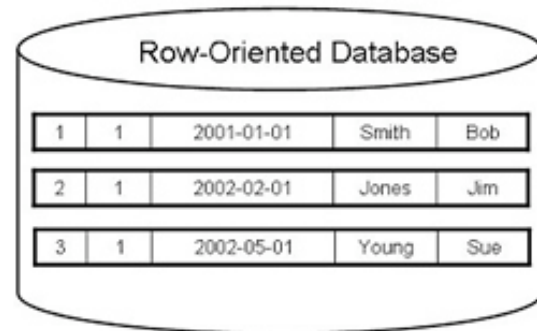
- Queries on keys
    - Scans on values

# Columnar (Relational) Database

(<https://www.dbbest.com/blog/column-oriented-database-technologies/>)

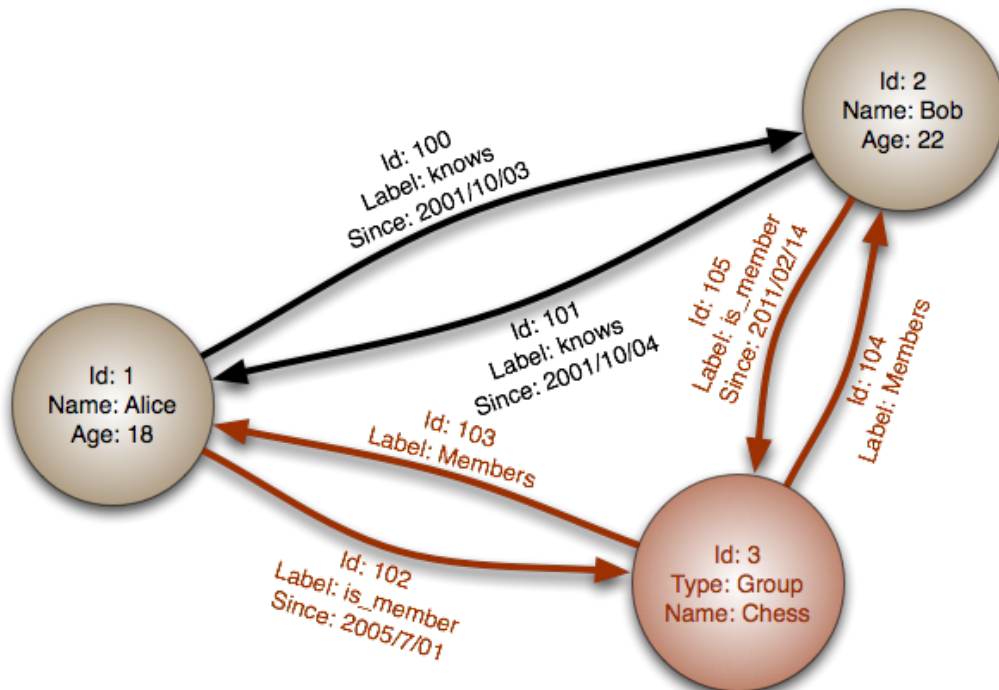
- Columnar and Row are both
  - Relational
  - Support SQL operations
- But differ in data storage
  - Row keeps row data together in blocks.
  - Columnar keys column data together in blocks.
- This determines performance for different types of query, e.g.
  - Columnar is extremely powerful for BI scenarios
    - Aggregation ops, e.g. SUM, AVG
    - PROJECT (do not load all of the row) to get a few columns
  - Row is powerful for OLTP. Transaction typically create and retrieve
    - One row at a time
    - All the columns of a single row.

Emp_no	Dept_id	Hire_date	Emp_in	Emp_fn
1	1	2001-01-01	Smith	Bob
2	1	2002-02-01	Jones	Jim
3	1	2002-05-01	Young	Sue
4	2	2003-02-01	Stemle	Bill
5	2	1999-06-15	Aurora	Jack
6	3	2000-08-15	Jung	Laura



# Graph Database

- Exactly what it sounds like
- Two core types
  - Node
  - Edge (link)
- Nodes and Edges have
  - Label(s) = “Kind”
  - Properties (free form)
- Query is of the form
  - $p1(n)-p2(e)-p3(m)$
  - $n, m$  are nodes;  $e$  is an edge
  - $p1, p2, p3$  are predicates on labels

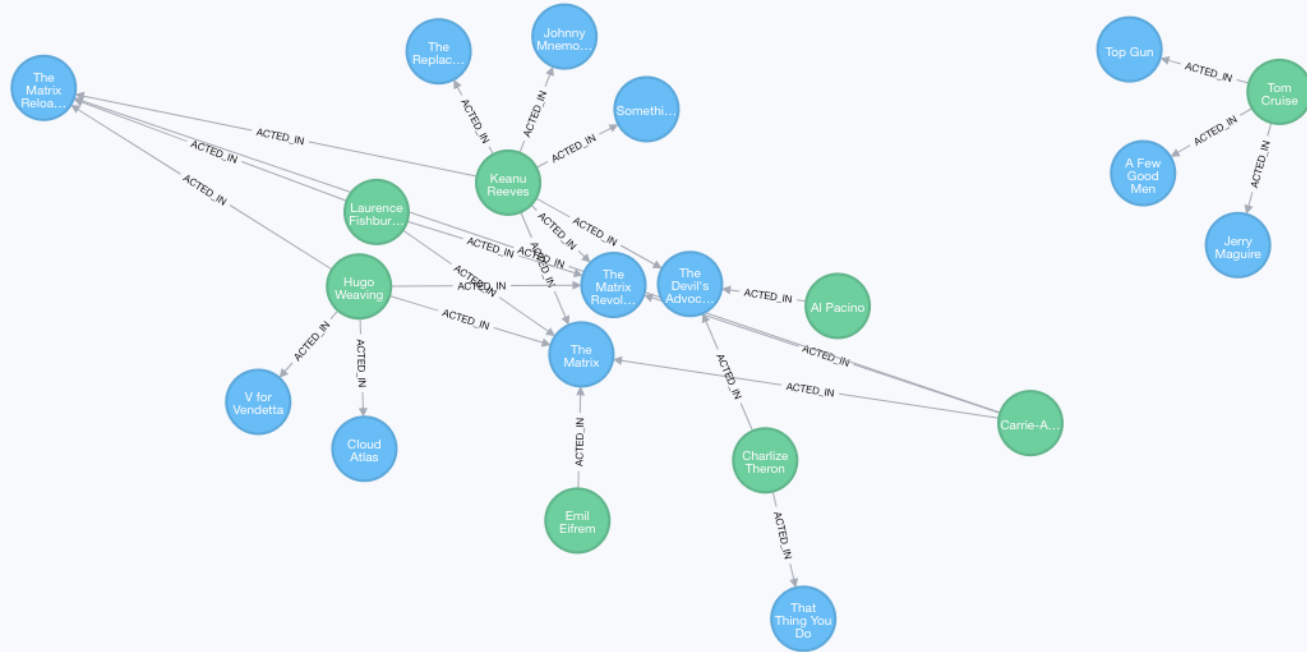


# Neo4J Graph Query

```
$ MATCH p=()-[r:ACTED_IN]->() RETURN p LIMIT 25
```

\*(21) Movie(13) Person(8)

\*(25) ACTED\_IN(25)



Displaying 21 nodes, 25 relationships.

AUTO-COMPLETE ☒

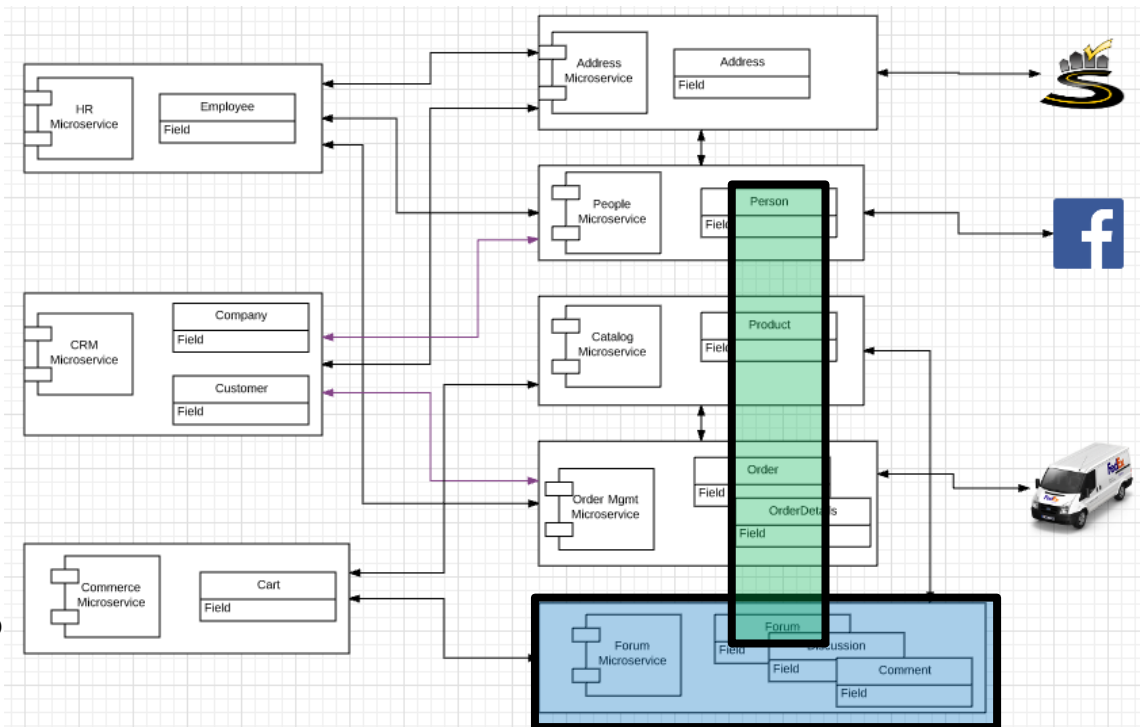
# Thanks, but Why?

The initial idea was

- DynamoDB for the forum
- Neo4J to track/query
  - Who bought what?
  - Who has bought things similar to whom?
  - Who commented on what?
  - etc.
- Use Redis to optimize
  - Idempotency
  - Etag

But

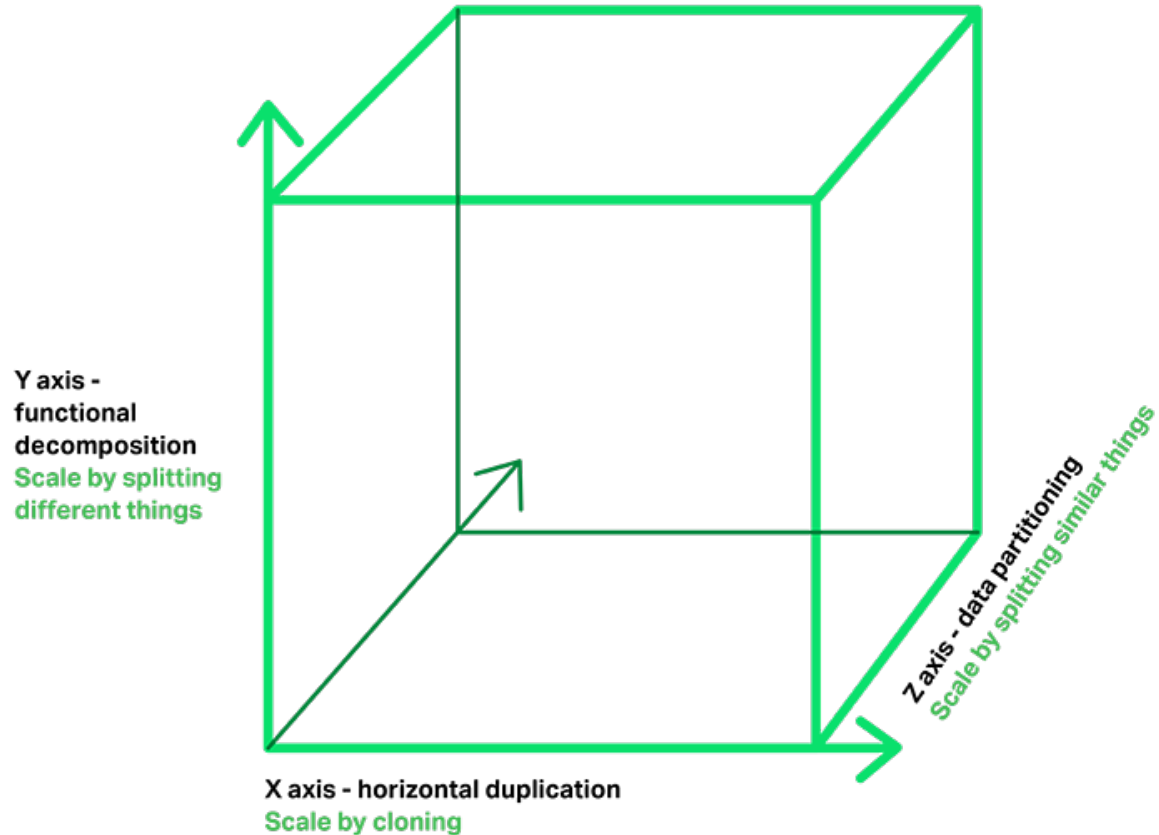
- we do not have enough time to do in context of solution.
- Will have to do smaller scenarios and use cases.





# Microservice XYZ Scaling

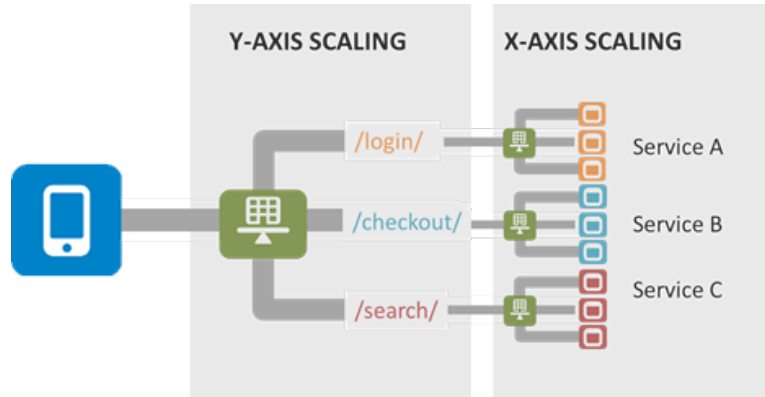
# Scaling Microservices



# Microservice Scaling

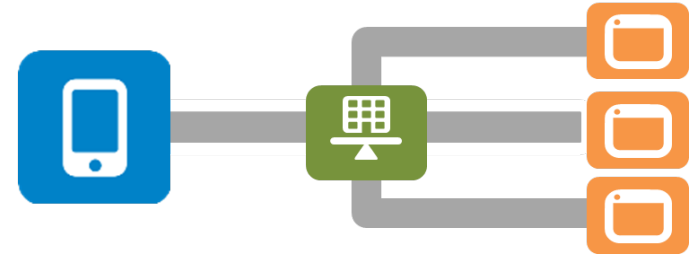
## Y-AXIS SCALING

Network name: Layer 7 Load Balancing, Content switching, HTTP Message Steering



## X-AXIS SCALING

Network name: Horizontal scaling, scale out



## Z-AXIS SCALING

Network name: Layer 7 Load Balancing, Content switching, HTTP Message Steering



<https://devcentral.f5.com/articles/the-art-of-scale-microservices-the-scale-cube-and-load-balancing>

# Microservice Scaling

## Y-AXIS SCALING

X-Axis (Horizontal Scaling) assumes

- Each instance has equal access to data →
- Highly scalable database in some scenarios.

And is the assumed model for

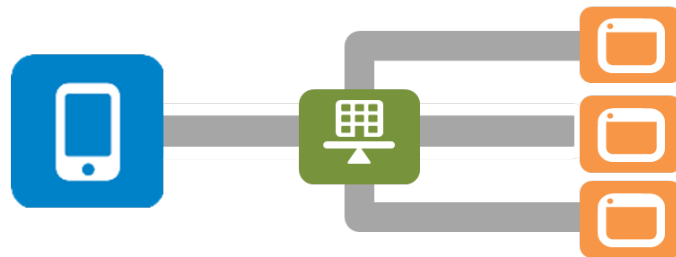
- Lambda functions
- Elastic BeanStalk

This is hard to achieve for scenarios like

- ETag
- Idempotency tokens

## X-AXIS SCALING

Network name: Horizontal scaling, scale out



## SCALING

Network name: Layer 7 Load Balancing, Content switching, HTTP Message Steering

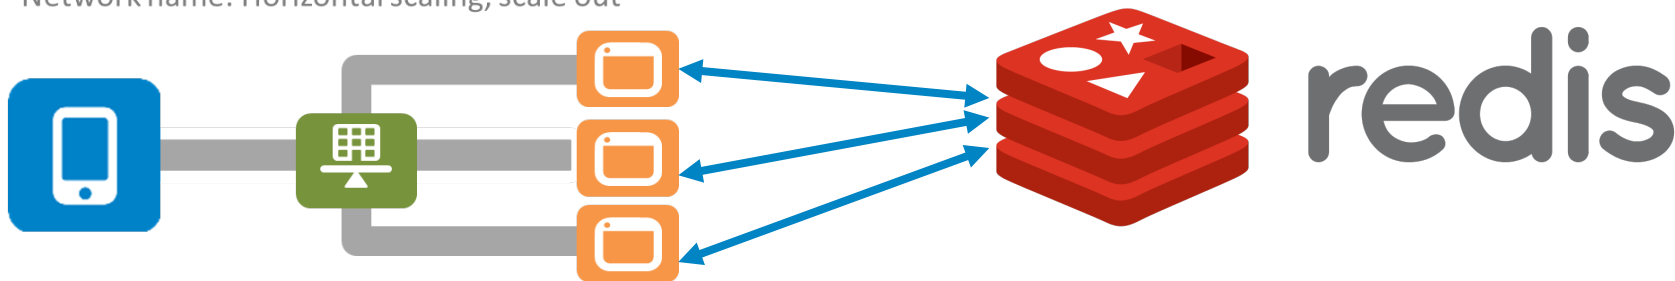


<https://devcentral.f5.com/articles/the-art-of-scale-microservices-the-scale-cube-and-load-balancing>

# Redis (Memcache), ETag, Idempotency

## X-AXIS SCALING

Network name: Horizontal scaling, scale out



- ETag
  - Store {URL, ETag} on GET cache miss.
  - Invalidate tag on any PUT, POST
  - Fail safe: If any doubt, e.g. partition, error, ..., treat as a conflict.
- Idempotency
  - Check for duplicate, and create if not exists.
  - Fail safe: do not process a request if any doubt.

# Similar to express-redis-cache

(<https://www.npmjs.com/package/express-redis-cache>)

Just use it as a middleware in the stack of the route you want to cache.

```
var app = express();
var cache = require('express-redis-cache')();

// replace
app.get('/',
  function (req, res) { ... });

// by
app.get('/',
  cache.route(),
  function (req, res) { ... });
```

This will check if there is a cache entry for this route. If not, it will cache it and serve the cache next time route is called.

# 12 Factor Applications

<https://12factor.net/>

<https://www.slideshare.net/bifer/twelve-factor-apps>

# The 12 Factors

## I. Codebase

One codebase tracked in revision control,  
many deploys

## II. Dependencies

Explicitly declare and isolate dependencies

## III. Config

Store config in the environment

## IV. Backing Services

Treat backing services as attached resources

## V. Build, release, run

Strictly separate build and run stages

## VI. Processes

Execute the app as one or more stateless processes

## VII. Port binding

Export services via port binding

## VIII. Concurrency

Scale out via the process model

## IX. Disposability

Maximize robustness with fast startup  
and graceful shutdown

## X. Dev/prod parity

Keep development, staging, and production  
as similar as possible

## XI. Logs

Treat logs as event streams

## XII. Admin processes

Run admin/management tasks as one-off processes



- See PDF