

COMS E6998: Microservices and Cloud Applications

Lecture 11: Q&A, Final Project, All Project Summary

Dr. Donald F. Ferguson
dff9@columbia.edu

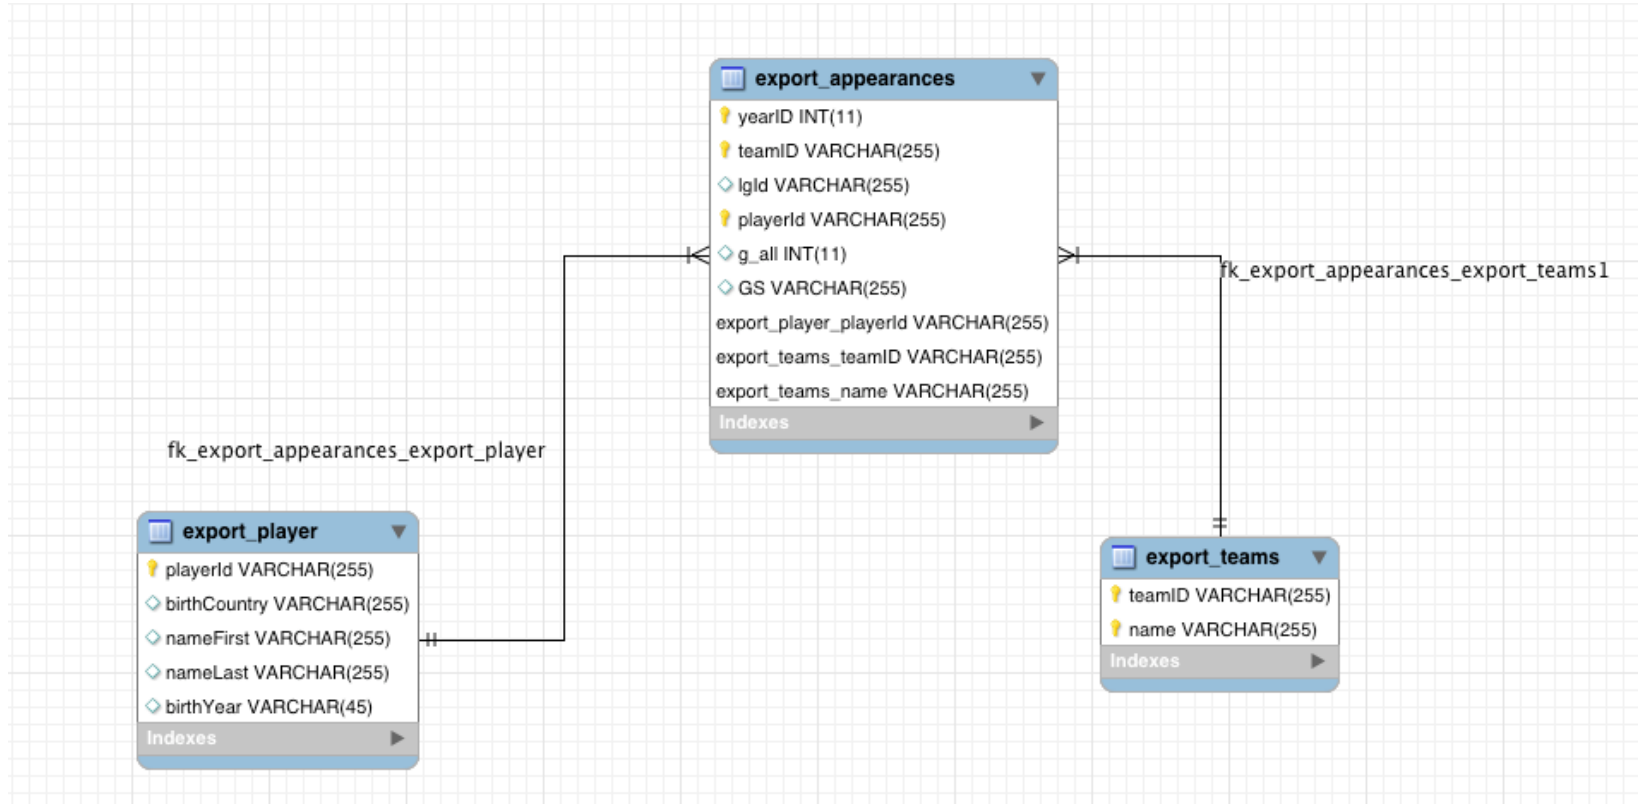
© Donald F. Ferguson, 2017. All rights reserved.

Comments Questions

Projects

Final Project

Logical Data Model



Standalone Neo4J Microservice

- Import data into Neo4J (local sever or in cloud)
 - Files are at <https://courseworks2.columbia.edu/courses/47421/files/#>
 - appearances.csv
 - teams.csv
 - players.csv
- Neo4J Import Process
 - Import from URL via HTTP or
 - From local directory (On Mac =
/Users/donaldferguson/Documents/Neo4j/default.graphdb/import)
 - Using LOAD CSV commands

Neo4J LOAD CSV

```
LOAD CSV WITH HEADERS FROM 'file:///teams.csv' AS line
CREATE (:Team { team_id: line.teamID, name: line.name})
```

```
LOAD CSV WITH HEADERS FROM 'file:///players.csv' AS line
CREATE (:Player { player_id: line.playerID, birthYear: line.birthYear,
                 birthCountry: line.birthCountry, first_name: line.nameFirst, last_name: line.nameLast})
```

USING PERIODIC COMMIT

```
LOAD CSV WITH HEADERS FROM "file:///appearances.csv" AS row
MATCH (p:Player {player_id: row.playerID}), (t:Team {team_id: row.teamID})
CREATE (p)-[:APPEARANCE {year: row.yearID, games: row.g_all}]->(t);
```

- Statement 1 creates nodes with label :Team for each row in CSV file
- Statement 2 creates nodes with label :Player for each row in CSV
- Statement 3: For each row in CSV
 - Find the node (p) :Player that has the row's playerID
 - Find the node (t) :Team that has the row's teamID
 - Create a relation :APPEARANCE from (p) to (t) with properties year, games to represent the fact that (p) *appeared* in *games* games in *year* year for team (t)

Find Player with player_id 'napolmi01'

\$ match (n:Player { player_id: 'napolmi01' }) return n


*(1) Player(1)

Graph

Table

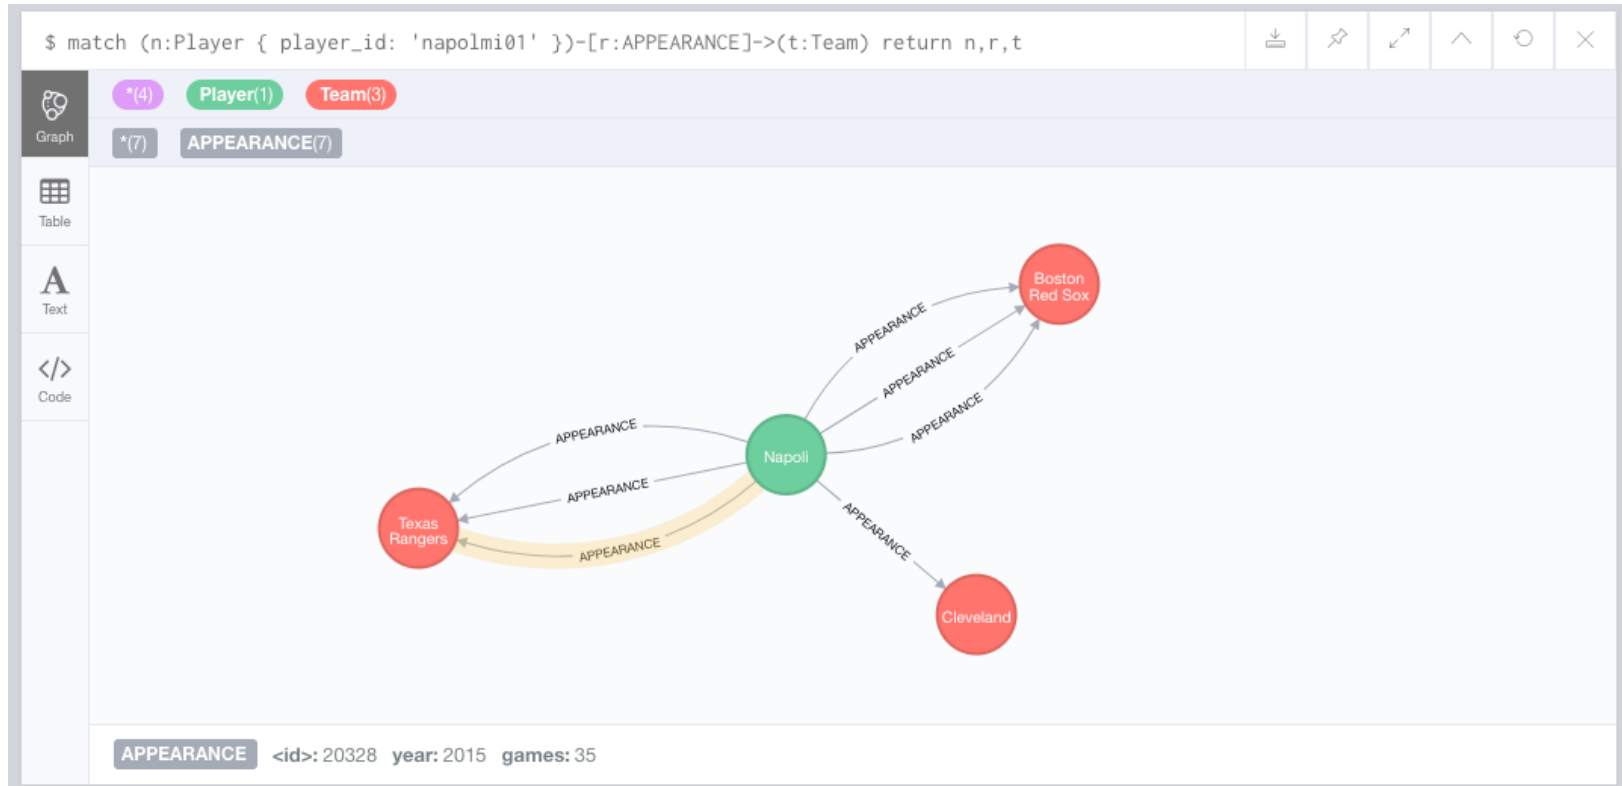
Text

Code



Player <id>: 301765 player_id: napolmi01 birthYear: 1981 birthCountry: USA last_name: Napoli first_name: Mike

Find the Teams/Years for Which Napoli Appeared




Final (Standalone) Project

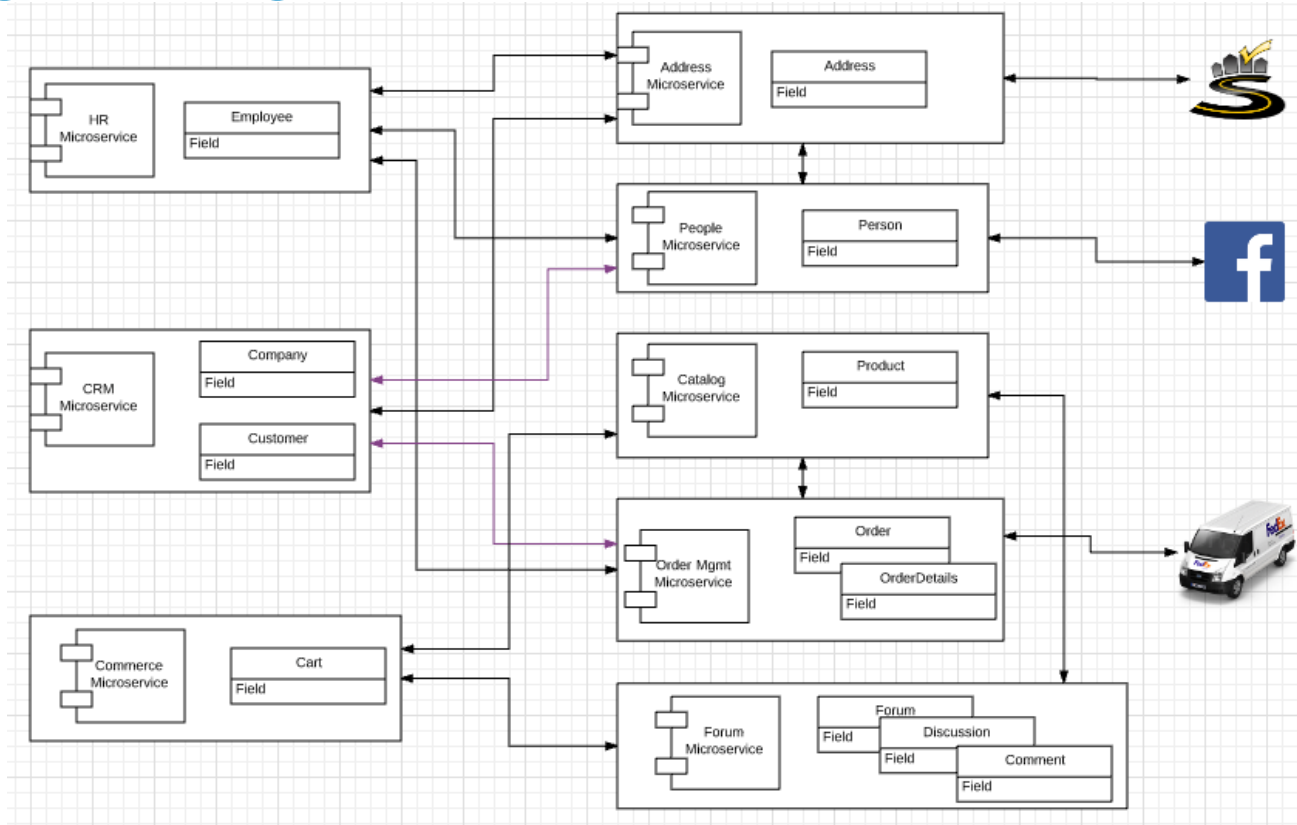
- Implement a Players microservice
 - /Players
 - GET (with query params)
 - POST
 - /Players/{playerid}
 - GET
 - PUT
 - DELETE
 - /Players/{id}/teams (GET only)
 - /Players/{id}/teammates
 - GET only.
 - One optional query parameter ?year=1234
 - Returns anyone with whom the player was on a team

This is Where it Gets Freaky

- Players A and B
 - Were teammates
 - If
 - A and B
 - Both appeared for
 - Any team T
 - In the same year.
- You will
 - Write queries to find
 - All of the teammates of
last_name='Napoli'
 - Using Neo4J and SQL



Original Big Picture



Other Projects

0th Project

- Teams
 - Form your teams (approx. 5 people)
 - Identify contact focal point.
 - Give your team a “cool” name.
- Signup (reuse) and Amazon Web Service Account
 - Free Tier should be fine.
 - Provide access to team members.
- Create an Elastic Beanstalk instance/application.
 - Use one of the sample application.
 - Will have to make more sophisticated starting next week;
I will use Node JS with Express for Elastic Beanstalk examples.

1st Project – Part 1

- Implement two distinct microservices
 - Person
 - Address
- Tasks
 - Use Swagger Editor to define and document REST APIs.
 - Implement an Elastic Beanstalk application (microservice) for each resource that implements the relevant REST API.
 - Each microservice should support
 - GET and POST on resource, e.g. /Person
 - GET, PUT, DELETE on resource/id, e.g. /Person/dff9
 - Simple query, e.g. /Person?lastName=Ferguson
 - Pagination
 - Relationship paths: /Person/dff9/address and /Addresses/someID/persons
 - HATEOAS links where appropriate.
 - Simple HTML/Angular demo UI.
- Due: 11:59 PM on 26-Sep-2017

Single Site

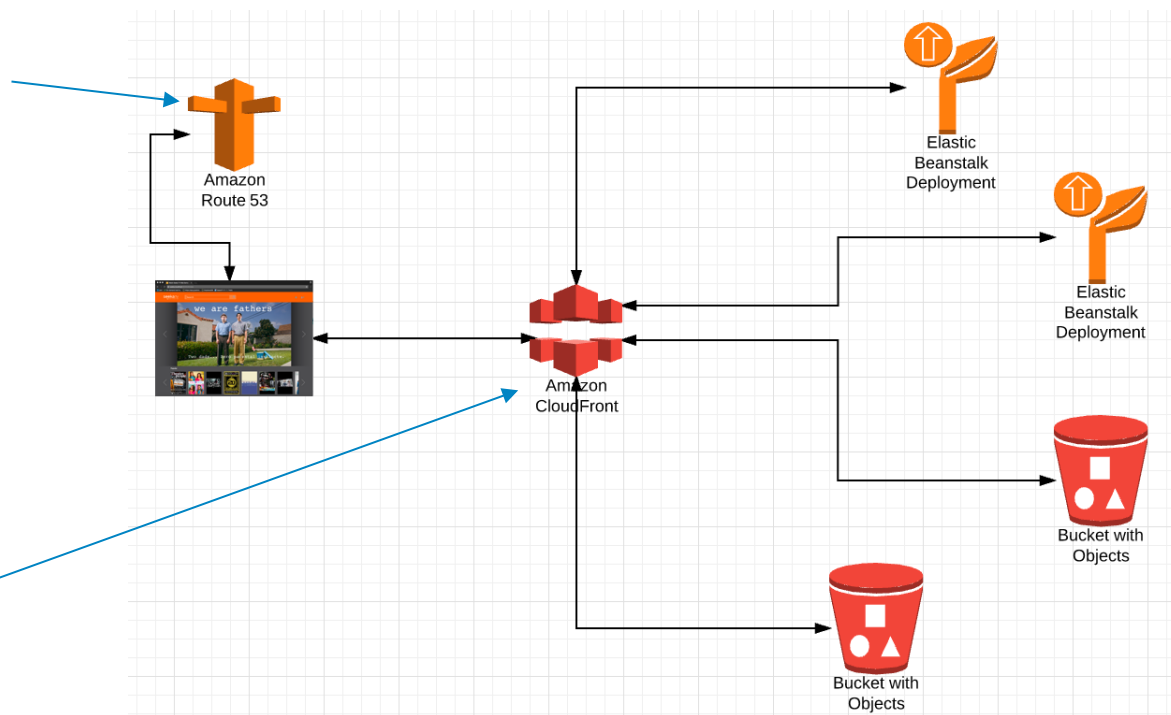
- DNS

- Resolve dff-company.com to
- Something.amazon.com
- Under the covers

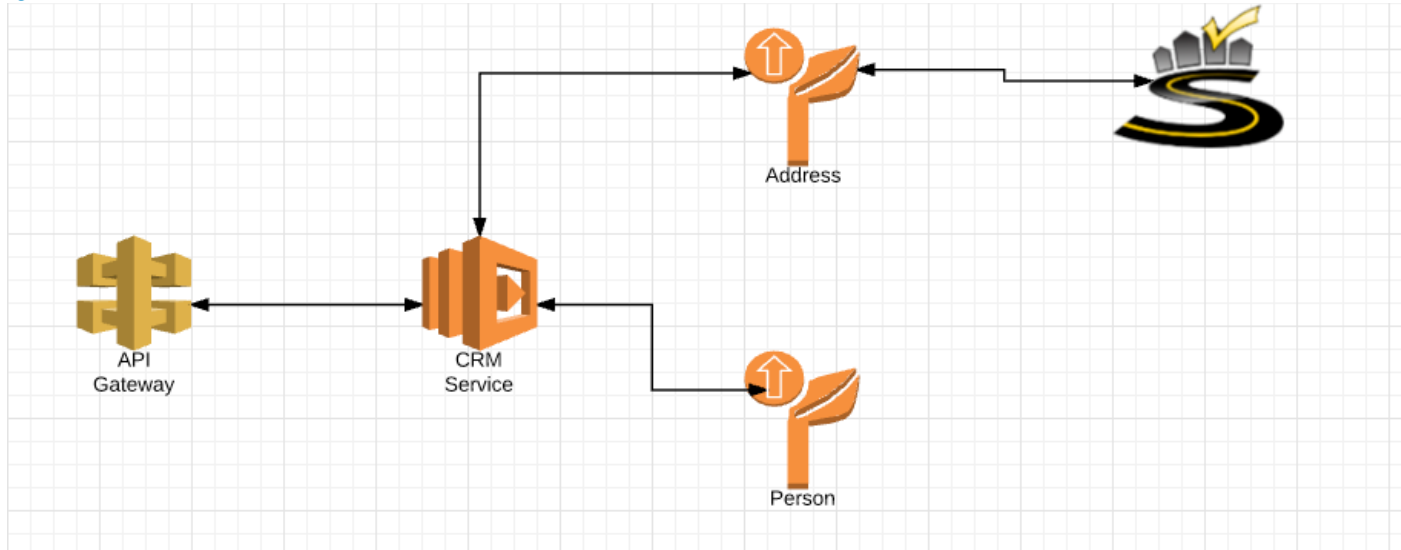
- Map

- /api/person
- /api/address
- /app
 - /js
 - /views
 - /app-content
- /digital-assets
 - /images
 - /videos

To correct IP addresses and sub-paths



Project 2 – Part 1



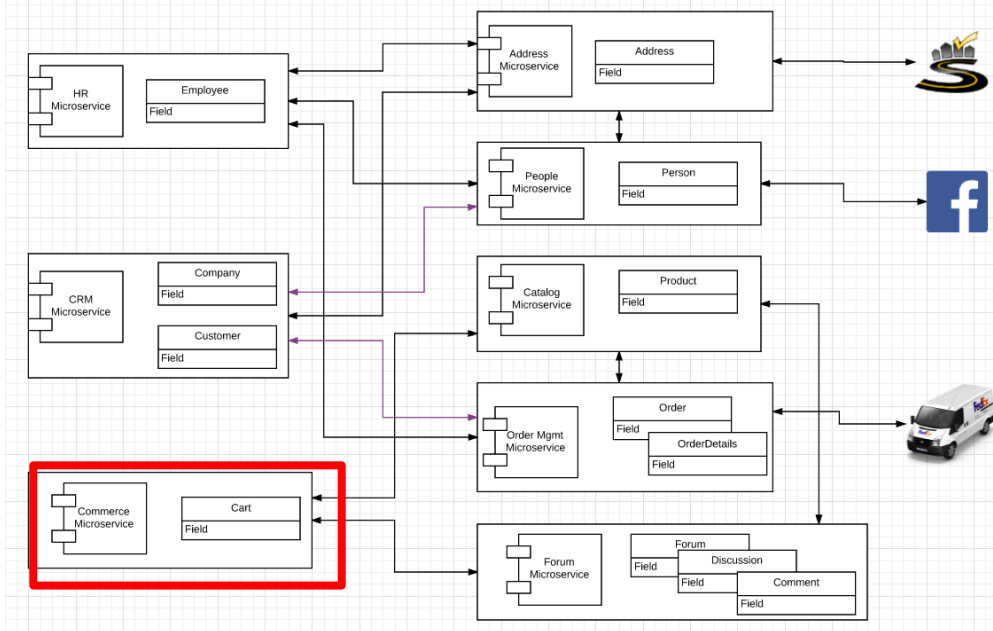
Build on 1st to microservices

- Implement the CRM Service (and HR Service) using Lambda functions, and orchestration approach from previous lecture.
- Integrate with SmartyStreets
- Deploy all 3 microservices via API Gateway.
- Deliver web content via CloudFront and S3.

Project 2 – Part 2

- Implement middleware plugins and microservices implementing
 - ETag generation and processing.
 - Idempotent functions.
 - SNS Event generation on PUT, POST and DELETE.
- Plug/deploy the plugin layer in
 - Each of the Beanstalk Microservices
 - The CRM Lambda function.
- Write an empty, placeholder Lambda function that reacts to the SNS events. We will do some interesting things with this later.

Project 3 4 – Start/Expand Commerce Microservice



- Use Lambda function

Customers may

- Register with
- And subsequently login with
- Facebook or Twitter

Just

- Implement register/login
- We will later
 - Use Step Functions for composing microservices and APIs
 - Publish commerce actions to FB and Twitter.

Also, write

- A placeholder
- API Gateway Custom Authorizer
- Which we will later use to manager authorization to orders.

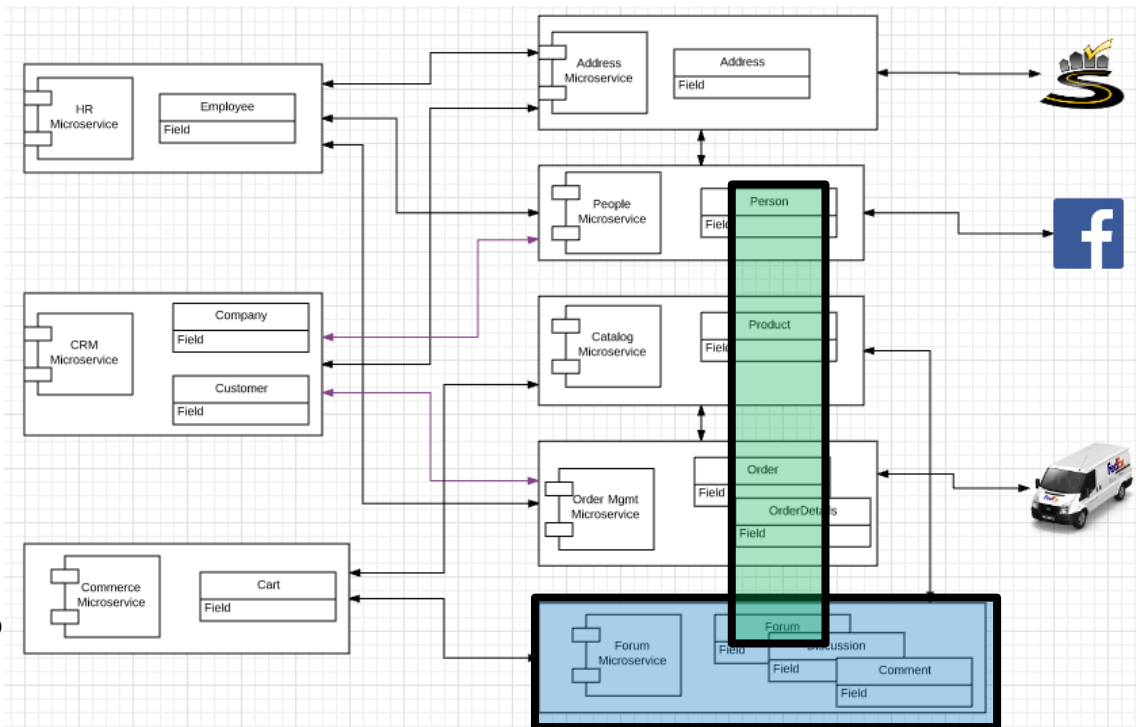
Last Week's Lecture – CAP and New Datamodels

The initial idea was

- DynamoDB for the forum
- Neo4J to track/query
 - Who bought what?
 - Who has bought things similar to whom?
 - Who commented on what?
 - etc.
- Use Redis to optimize
 - Idempotency
 - Etag

But

- we do not have enough time to do in context of solution.
- Will have to do smaller scenarios and use cases.



My View on Status

Project	Subtask	Element
Project 0	Set up team Run ELB	
Project 1	Implement 2 microservices	<ul style="list-style-type: none"> • Implement • REST API/HATEOAS • Query, relationships
	Swagger	<ul style="list-style-type: none"> • Document, test • Integrate with API GW, development tools, ...
Project 2	Cloud APIs	<ul style="list-style-type: none"> • SmartyStreets (UI, backend)
	Orchestration	<ul style="list-style-type: none"> • In code
	Single Site Image	<ul style="list-style-type: none"> • S3 for content • CloudFront • API Gateway
	Middleware	<ul style="list-style-type: none"> • ETag and idempotency technical microservice • SNS event generation
Project 3	OAuth2	<ul style="list-style-type: none"> • Logon, Register • Custom API GW Authorizer → Authorization
Additional Goals		<ul style="list-style-type: none"> • Step functions • SNS to integrate registration, commerce views, etc. with Neo4J • SNS → Swagger via WebHooks/Lambda • Redis: Optimize data sharing for middleware scenarios • Rules engines • Text search

Timetable

- Project 2 and project 3
 - Complete (initial) reviews by 05-Dec.
 - Use 27-Nove and 05-Dec lectures if possible, then office hours, then extra OH.
- We will hold one final project review with each team
 - Covering all elements of projects 1, 2 and 3
 - Including modification/correction based on feedback from prior reviews.
 - I will provide a list of completion requirements metrics.
 - Project 4, which has two small subelements
 - Simple idempotency check microservice → Redis
 - POST token returns OK if token not previously seen.
 - POST returns error if token previously recorded.
 - Simple graph microservice on Neo4J using sample data.
 - POST a simple query command in JSON.
 - Returns results from Neo4J.